

Deep Learning Series

Episode 3 - “La revanche des token”

Recap from last episode


What is Deep Learning ?

Deep Learning is a **subfield of Machine Learning** : A specific way of learning representations from data that puts an emphasis on **learning successive layers of increasingly meaningful representations**

Other approaches to **Machine Learning** tend to focus on learning only one or two layers of representations of the data -> **Shallow Learning**


Deep Learning = Deep Sequence of simple data transformations (layers)

Simple Neural Network



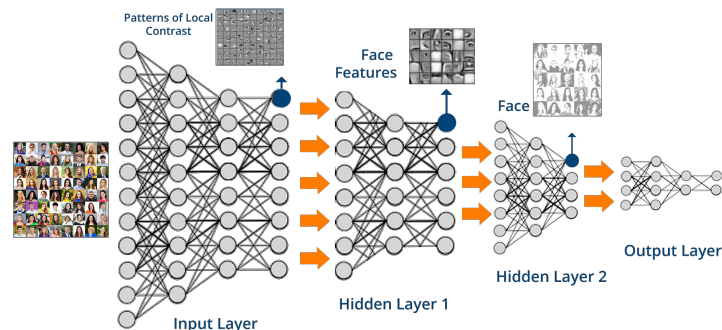
● Input Layer

Deep Learning Neural Network



● Hidden Layer


● Output Layer



*Layered representations of the data are almost always learned via models called **Neural Networks***

How Deep Learning works ?

The Deep Learning workflow



A **loss function** measures the quality of the network's output.

It computes a distance score between the prediction and the true target

Fighting overfitting : Reducing the network's size


Simplest way to reduce overfitting : Reducing the number of learnable parameters (=degrees of freedom) in the model

- A model with more parameters has more **memorization capacity**
- With fewer learnable parameters, the model will have to learn **compressed representations** to minimize its loss
- Find the right amount of parameters by tuning the model's configuration

```
original_model = models.Sequential()  
original_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
original_model.add(layers.Dense(16, activation='relu'))  
original_model.add(layers.Dense(1, activation='sigmoid'))
```



```
smaller_model = models.Sequential()  
smaller_model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))  
smaller_model.add(layers.Dense(4, activation='relu'))  
smaller_model.add(layers.Dense(1, activation='sigmoid'))
```



Fighting overfitting : Adding weight regularization


Principle of Occam's razor : Given two explanations for something, the explanation most likely to be correct is the simplest one.

- Simpler models are less likely to overfit than complex ones
- Simpler model => Forcing its weights to take only small values (more regular distribution of weights)
- This is done by adding to the loss function a cost associated with having large weights (L1 or L2 regularization)

```
original_model = models.Sequential()  
original_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
original_model.add(layers.Dense(16, activation='relu'))  
original_model.add(layers.Dense(1, activation='sigmoid'))
```




```
l2_model = models.Sequential()  
l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),  
                           activation='relu', input_shape=(10000,)))  
l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),  
                           activation='relu'))  
l2_model.add(layers.Dense(1, activation='sigmoid'))
```



Fighting overfitting : Adding dropout

Dropout : Randomly dropping out (setting to 0) a number of output features of a given layer during training. The dropout rate is usually between 0.2 and 0.5. Use all outputs at test time.


- Core idea : Introducing noise in the output values can **break up “bad luck circumstance patterns”** in the training data. Information need to be encoded at several places to be relevant.



```
original_model = models.Sequential()  
original_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
original_model.add(layers.Dense(16, activation='relu'))  
original_model.add(layers.Dense(1, activation='sigmoid'))
```



```
dpt_model = models.Sequential()  
dpt_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
dpt_model.add(layers.Dropout(0.5))  
dpt_model.add(layers.Dense(16, activation='relu'))  
dpt_model.add(layers.Dropout(0.5))  
dpt_model.add(layers.Dense(1, activation='sigmoid'))
```

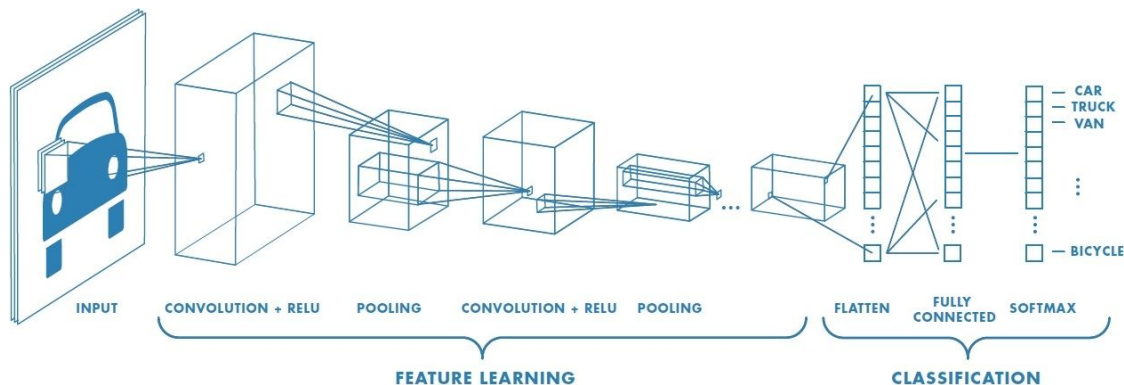


What did we see ?

Understanding Convolutional Neural Networks

Using Data Augmentation

Using pretrained networks



| Pretrained Network for Feature Extraction

- **Convolutional Neural Networks** are the best type of machine learning models for computer vision tasks
 - It is possible to train them from scratch even on very small datasets
- On a small dataset, overfitting will be the main issue. **Data augmentation** is a powerful way to fight it for image data
- It's easy to reuse an existing convnet that was trained on a larger dataset for **Feature Extraction**
- **Fine-tuning** adapts to a new problem some of the representations previously learned by an existing model
 - This pushes performances a bit further

Deep Learning for text and sequences

Embeddings & Recurrent Neural Networks


| What will we see ?

Embeddings for text data


Recurrent Neural Networks

1D Convolutional Neural Networks


one to one




one to many




many to one



many to many



many to many



Source :

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

I Quelles applications ?

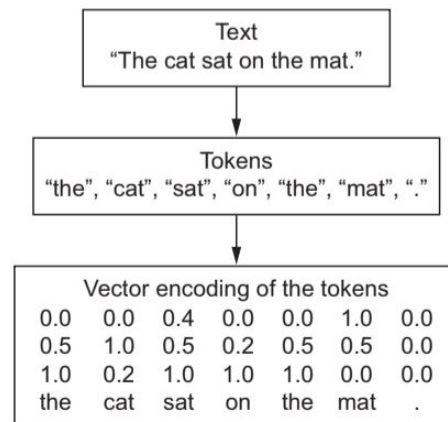
- Document and time series **classification**
 - *Ex : Identifying the topic of an article, the author of a book*
- Time series comparisons
- **Sequence-to-Sequence** learning
 - *Ex : Decoding an English sentence into French*
- Sentiment Analysis
- Time series **forecasting**

Working with Text Data

Lorem ipsum dolor sit amet, consectetur adipiscing elit.


Vectorizing text

- Deep Learning models don't take raw text as input, they only work with **numerical tensors**
 - We need to convert sentences into sequences of vectors
 - 1 **token** (1 word) = 1 vector
- How to associate a vector with a token ?
 - **One-Hot encoding** : Sparse vectors
 - **Token Embedding** : Dense, shorter vectors



| One-Hot encoding

- One of the most common way to turn a token into a vector
- How to create the vectors ?
 - Associate a **unique integer index** with every word (N words)
 - Turn this integer index i into a **binary vector** of size N , which is all-zeros except for the i th entry, which is 1
- Variant : **One-Hot hashing trick**
 - When your vocabulary size is too large
 - Hash words into vectors of fixed size
 - But susceptible to *hash collisions*




Word Embeddings

- Word Embeddings are low-dimensional floating-point vectors
 - Common sizes : 256, 512, 1024
- They are learned from data
- Two ways to obtain word embeddings :
 - Learn embeddings **jointly with the main task** you care about
 - Start with random word vectors and learn them through back-propagation
 - Load a **pre-computed word embeddings** learned using a different machine learning task
 - **Pre-trained word embeddings**



One-hot word vectors:

- Sparse
- High-dimensional
- Hardcoded




Word embeddings:

- Dense
- Lower-dimensional
- Learned from data

How to interpret word embeddings ?

- The geometric relationships between word vectors should reflect the **semantic relationships** between these words
 - Word embeddings are meant to map human language into a geometric space
- We would expect the geometric distance between two semantically related words to be short
- We also want specific directions in the embedding space to be meaningful
 - Ex : “Female vector” to convert king into queen



From Mikolov et al. (2013a)

$$V(\text{King}) - V(\text{Man}) + V(\text{Woman}) \approx V(\text{Queen})$$

Learning word embeddings with the Embedding Layer

- There is **no ideal word embedding space** that would perfectly map human language and could be used for any task
- Need to **learn a new embedding space** with every new task
- There is a special layer for that in Keras : the **Embedding** layer

```
from keras.layers import Embedding

# The Embedding layer takes at least two arguments:
# the number of possible tokens, here 1000 (1 + maximum word index),
# and the dimensionality of the embeddings, here 64.
embedding_layer = Embedding(1000, 64)
```

- The *Embedding* layer can be seen as a dictionary that maps integer indices to dense vectors
 - It takes integers as input
 - Returns the associated vector
- **Input** : 2D tensor of shape (*samples*, *sequence_length*)
- **Output** : 3D tensor of shape (*samples*, *sequence_length*, *embedding_dimensionality*)

Learning word embeddings with the Embedding Layer

➤ Example for sentiment prediction

```
from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()
# We specify the maximum input length to our Embedding layer so we can later flatten the embedded inputs
model.add(Embedding(10000, 8, input_length=maxlen))
# After the Embedding layer, our activations have shape `(samples, maxlen, 8)`.
# We flatten the 3D tensor of embeddings into a 2D tensor of shape `(samples, maxlen * 8)`
model.add(Flatten())
# We add the classifier on top
model.add(Dense(1, activation='sigmoid'))
```


Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 20, 8)	80000
flatten_1 (Flatten)	(None, 160)	0
dense_1 (Dense)	(None, 1)	161



By flattening the embedded sequences, the model treats each word in the input sequences separately, without considering **inter-word relationships** and **sentence structure**.

Using pre-trained word embeddings

- When not enough data is available, it is hard to learn meaningful word representations from scratch
- Instead of learning word embeddings jointly with the main task, you can load embedding vectors from a **pre-computed embedding space**
 - Which is highly structured
 - That capture generic aspects of language structure
- It is the same approach as using pre-trained convnets in image classification
- Such embeddings are generally computed using **word-occurrence statistics**
 - They are trained with text corpus like Wikipedia



Using pre-trained word embeddings

➤ Example for sentiment prediction

```
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
```

- The *Embedding* layer has a single weight matrix : a 2D float matrix
 - Each entry i is the word vector associated with index i

```
model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False
```

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 100, 100)	1000000
flatten_2 (Flatten)	(None, 10000)	0
dense_2 (Dense)	(None, 32)	320032
dense_3 (Dense)	(None, 1)	33
Total params: 1,320,065		
Trainable params: 320,065		
Non-trainable params: 1,000,000		



Don't forget to precise the first layer to be **not trainable**.

The pre-trained parts should not be updated during training, to avoid forgetting what we already know.


Understanding Recurrent Neural Networks

Don't forget the past

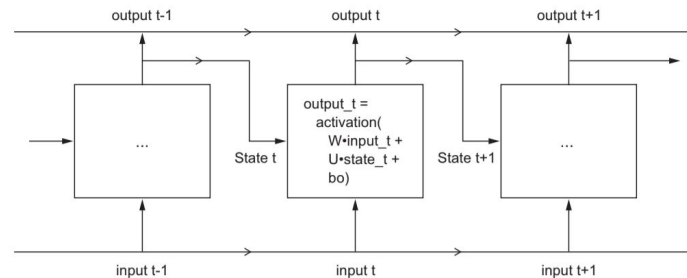
Intuition

- Densely connected networks and convnets have **no memory** : each input is processed independently
 - To process sequence or time series data, you need to show the entire sequence to the network at once
- Humans process a sequence word by word, while keeping memories of what came before
- A **Recurrent Neural Network** adopts the same principle, in a very simplified version
 - Processes a sequence by iterating through it and maintaining a state containing past information
 - It's a neural network with an **internal loop**

```
state_t = 0
for input_t in input_sequence:
    output_t = activation(dot(W, input_t) + dot(U, state_t) + b)
    state_t = output_t
```



Recurrent Neural Network : A network with a loop



A simple RNN, unrolled over time

| Recurrent Layer in Keras

- This simple process is present in Keras as a **SimpleRNN** layer

```
from keras.layers import Dense

model = Sequential()
model.add(Embedding(max_features, 32))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))
```


- This layer can be run in two modes
 - Return the full sequences of successive outputs for each timestep
 - Return only the last output for each input sequence
- These two modes are controlled by the **return_sequences** constructor argument



If you stack several recurrent layers one after the other in order to increase the representational power of the network, all intermediate layers must return full sequences of outputs.

Understanding the LSTM and GRU layers

- In practice, one never uses the **SimpleRNN** layer
 - They are too simplistic to be of real use
- Major issue : it has high difficulties to learn long-term dependencies
 - Due to the **vanishing gradient problem**
- **LSTM** and **GRU** layers are designed to solve this problem
 - LSTM : Long Short-Term Memory
 - GRU : Gated Recurrent Unit
- LSTM
 - Adds a way to carry information across many timesteps
 - Information from a timestep can be transported into another timestep
- GRU
 - Use the same principle as LSTM
 - Cheaper to run but less representational power



The repeating module in an LSTM contains four interacting layers.

```
from keras.layers import LSTM

model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
```

Advanced use of RNNs

Towards infinity, and beyond

| Recurrent dropout to fight overfitting

- Like with densely connected and convolutional neural networks, **dropout** can be used for Recurrent Neural Networks to reduce **overfitting**
 - The **same dropout mask** should be applied at every timestep
 - Allows to properly propagate its learning error through time
- In Keras, there are two dropout-related arguments
 - **dropout** : dropout rate for input units
 - **recurrent_dropout** : dropout rate of the recurrent units

```
from keras.models import Sequential
from keras import layers
from keras.optimizers import RMSprop

model = Sequential()
model.add(layers.GRU(32,
                    dropout=0.2,
                    recurrent_dropout=0.2,
                    input_shape=(None, float_data.shape[-1])))
model.add(layers.Dense(1))
```

Stacking recurrent layers

- **Stacking** recurrent layers helps increasing the **representational capacity of the network**
 - What currently power the Google Translate algorithm is a stack of seven large LSTM layers
- To stack recurrent layers on top of each other in Keras, all intermediate layers should return their full sequence of outputs rather than their output at the last timestep
 - This is done by specifying *return_sequences=True*

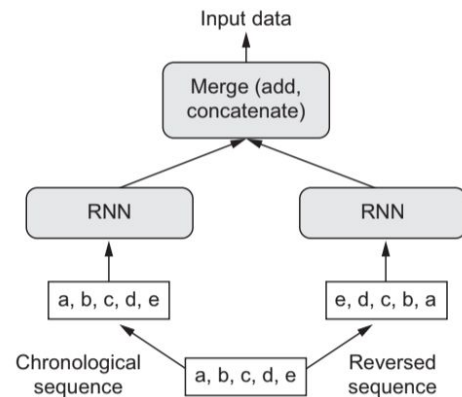
```
from keras.models import Sequential
from keras import layers
from keras.optimizers import RMSprop

model = Sequential()
model.add(layers.GRU(32,
                    dropout=0.1,
                    recurrent_dropout=0.5,
                    return_sequences=True,
                    input_shape=(None, float_data.shape[-1])))
model.add(layers.GRU(64, activation='relu',
                    dropout=0.1,
                    recurrent_dropout=0.5))
model.add(layers.Dense(1))

model.compile(optimizer=RMSprop(), loss='mae')
```

Bidirectional Recurrent Neural Networks

- Variant of the common RNN that is frequently used in natural-language processing
- It exploits the order-sensitivity of RNNs
- Consists of using **two regular RNNs** (LSTM or GRU), each of which processes the input sequence in one direction
 - Chronologically and anti-chronologically
 - Then merges their representations
- By processing a sequence both ways, it can catch patterns that may be overlooked by an unidirectional RNN



```
model = Sequential()  
model.add(layers.Embedding(max_features, 32))  
model.add(layers.Bidirectional(layers.LSTM(32)))  
model.add(layers.Dense(1, activation='sigmoid'))  
  
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])  
  
history = model.fit(x_train, y_train, epochs=10, batch_size=128, validation_split=0.2)
```


Sequence processing with convnets

What ? Convnets ? For sequence data ??

Understanding 1D convolution for sequence data


- Convolutional Neural Networks perform particularly well on image data
 - They are composed of 2D convolutions
- The same principle can be applied for sequence data with **1D convolutions** to extract local **1D patches** from sequences
- Such layers can recognize local patterns in a sequence
 - Translation invariant since the same input transformation is performed on every patch
- 2D pooling was used to spatially downsample image tensors
 - We can use the 1D equivalent : **1D pooling**

```
model = Sequential()  
model.add(layers.Embedding(max_features, 128, input_length=max_len))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.MaxPooling1D(5))  
model.add(layers.Conv1D(32, 7, activation='relu'))  
model.add(layers.GlobalMaxPooling1D())  
model.add(layers.Dense(1))
```



Combining CNNs and RNNs to process long sequences

- 1D convnets process input patches independently
 - They aren't sensitive to the order of timesteps
 - It often fails at giving meaningful results in cases when the order counts
 - But it is cheaper to train than RNNs
- RNNs have other characteristics
 - Longer to train, but can catch temporal dependencies
 - Still difficult to train on very long sentences
- One solution : **Combining CNNs and RNNs**
 - Use a 1D-convnet as a preprocessing step before an RNN
 - The convnet will downsample the sequence



Take aways

What did we learn ?

| What did we learn ?

- Techniques for working with sequence data
 - Tokenize text
 - **Embeddings**
 - RNNs : **SimpleRNN, LSTM, GRU**
- Increase performances of RNNs by using dropout, stacking layers and using bidirectional RNNs
- Use **1D convnets** and combine them with RNNs
- If global order matters (ex: time series), it's preferable to use a RNN
- If global order isn't fundamentally meaningful, using a 1D convnet are cheaper and work just as well